

```
import random
import numpy as np
from tqdm import tqdm
```

## ✓ Rover

### Monte Carlo Control



Consider a rover that is trying to move up the hill. The robot has solar powered motors. Once the robot **reaches the top** of the hill, the **episode ends**. The Robot can either be at **low, medium, high or top** of the hill at any given point of time.

The robot can either decide to spin it's wheels **slowly or rapidly** in order to move up on the gradient.

- If the **motor spins the wheel slowly**, with probability  $0.6$  it moves to the next higher state in one time step, and with probability  $0.4$ , it slides all the way down the slope to the low state.
- On the other hand, if the **motor spins the wheel rapidly**, with probability  $0.9$  it moves to the next higher state in one time step, and with probability  $0.1$ , it slides all the way down the slope to the low state.

The robot initially starts from low state with 10 units of energy (battery charge). The **episode ends** if the batteries get completely discharged or if rover reaches the top position.

The rover has a motor that can spin its wheel

- slowly at the expense of 2 unit of energy per time step;
- or rapidly at the expense of 4 units of energy per time step.

**rover gains 4 units of energy** per time step from its solar panels any time it transitions **upward from the medium position** as it gets exposed to sunlight.

Goal:

The Robot has to reach the top with maximum amount of energy left in the battery (Energy spent while moving to the top must be minimum)

## ✓ Tasks

### PART 1

- Create an "Environment" class that has the step() method which takes in an action and returns the next state and one step reward.
- Try taking random actions until the episode terminates. Print the one-step reward, action and next state for all transitions in the episode.  
What is the average total reward if you run the experiment over many episodes?
- What is the average total reward over 10,000 episodes if:
  - Actions are chosen "randmoly" from the action space?
  - Action is always "spin\_slowly"?
  - Acion is always "spin\_rapidly"?

### PART 2

- Create an "Agent" class with train() and policy() methods which learns to take actions in this environment. Agent should be trainable using Monte-Carlo Control.
- Train the agent. What is the average total return after training the Agent?  
(optional)
- Perform monte-carlo prediction of state-values to evaluate the learned policy. How valuable is it to start from each state?

```
INITIAL_CHARGE = 10
INITIAL_STATE = "low"
```

```
class RoverEnv:
```

```
    """It can be inferred from the problem that:
```

```
    State space = {"low", "medium", "high", "top"}
```

```
    Action space = {spin_slowly, spin_rapidly}
```

```
    Ex. `action = spin_slowly` means the agent took the action spin_slowly.
```

```
    Start state
```

```
    -----
```

```
    Robot starts low position with 10 units charge.
```

```
    Rewards
```

```
    -----
```

```
    Amount of energy left in the robot when episode ends is the reward.
```

```
    the one step-reward is received only once i.e, when the episode ends.
```

```
    Terminal conditions
```

```
    -----
```

```
    Episode end is:
```

```
    - Robot transitions to the top position.
```

```
    - There is no charge left in the batteries.
```

```
    """
```

```
    def __init__(self):
```

```
        self.actions = ["spin_slowly", "spin_rapidly"]
```

```
        self.action_space = self.actions.copy()
```

```
        #self.action_space = ["spin_slowly", "spin_rapidly"]
```

```
        self.states = ["low", "medium", "high", "top"]
```

```
        self.observation_space = self.states.copy()
```

```
        #self.observation_space = ["low", "medium", "high", "top"]
```

```
        self.reset()
```

```
    def reset(self):
```

```
        self.state = INITIAL_STATE
```

```
        self.charge = INITIAL_CHARGE
```

```
        self.terminated = False
```

```
        self.total_reward = 0
```



```
def get_transition_probability(self, start_state, end_state, action):
```

```
    # probability of moving "up" when motor spins "slowly"
    slow_spin_up_p = 0.6
    # probability of moving "up" when motor spins "rapidly"
    rapid_spin_up_p = 0.9

    # probability of moving "down" when motor spins "slowly"
    slow_spin_down_p = 1-slow_spin_up_p
    # probability of moving "down" when motor spins "rapidly"
    rapid_spin_down_p = 1-rapid_spin_up_p
```

```
    if action == "spin_slowly":
        if start_state == "low" and end_state == "medium":
            return slow_spin_up_p
        if start_state == "low" and end_state == "low":
            return slow_spin_down_p
        if start_state == "medium" and end_state == "high":
            return slow_spin_up_p
        if start_state == "medium" and end_state == "low":
            return slow_spin_down_p
        if start_state == "high" and end_state == "top":
            return slow_spin_up_p
        if start_state == "high" and end_state == "low":
            return slow_spin_down_p
```

```
    if action == "spin_rapidly":
        if start_state == "low" and end_state == "medium":
            return rapid_spin_up_p
        if start_state == "low" and end_state == "low":
            return rapid_spin_down_p
        if start_state == "medium" and end_state == "high":
            return rapid_spin_up_p
        if start_state == "medium" and end_state == "low":
            return rapid_spin_down_p
        if start_state == "high" and end_state == "top":
            return rapid_spin_up_p
        if start_state == "high" and end_state == "low":
            return rapid_spin_down_p
```

```
    return 0
```

```
def get_charge_difference(self, start_state, end_state, action):
```



```
charge_diff = 0

if action == "spin_slowly":
    charge_diff -= 2
elif action == "spin_rapidly":
    charge_diff -= 4

if start_state == "medium" and end_state == "high":
    charge_diff += 4
elif start_state == "medium" and end_state == "high":
    charge_diff += 4

return charge_diff

def get_reward(self, end_state):
    if end_state == "top":
        return self.charge
    return 0

def step(self, action):

    if self.state == "top":
        self.terminated = True
        return self.state, self.charge, self.terminated

    next_states = self.observation_space
    transition_probs = [self.get_transition_probability(self.state, next_state, action) for next_state in next_states]
    next_state = np.random.choice(next_states, p=transition_probs)

    self.charge += self.get_charge_difference(self.state, next_state, action)
    one_step_reward = self.get_reward(next_state)

    self.state = next_state

    if self.state == "top" or self.charge <=0:
        self.terminated = True
```

```
self.total_reward+=one_step_reward
```

```
return self.state, one_step_reward, self.terminated
```

```
class Agent:
```

```
    def __init__(self, env):
```

```
        self.env = env
```

```
        self.Q = {state: {action: 0 for action in env.actions} for state in env.states}
```

```
        self.returns = {state: {action: [] for action in env.actions} for state in env.states}
```

```
        self.policy = {state: random.choice(env.actions) for state in env.states}
```

```
    def train(self, num_episodes=10000, epsilon=0.1):
```

```
        for _ in range(num_episodes):
```

```
            self.env.reset()
```

```
            state = self.env.state
```

```
            episode = []
```

```
            while not self.env.terminated:
```

```
                action = self.policy[state] if random.random() > epsilon else random.choice(self.env.actions)
```

```
                next_state, reward, _ = self.env.step(action)
```

```
                episode.append((state, action, reward))
```

```
                state = next_state
```

```
            G = 0
```

```
            for state, action, reward in reversed(episode):
```

```
                G = reward + G
```

```
                self.returns[state][action].append(G)
```

```
                self.Q[state][action] = np.mean(self.returns[state][action])
```

```
                self.policy[state] = max(self.Q[state], key=self.Q[state].get)
```

```
    def get_action(self, state):
```

```
        return self.policy[state]
```

```
    def reset(self):
```

```
        self.Q = {state: {action: 0 for action in env.actions} for state in env.states}
```

```
        self.returns = {state: {action: [] for action in env.actions} for state in env.states}
```

```
        self.policy = {state: random.choice(env.actions) for state in env.states}
```



```
env = RoverEnv()
agent = Agent(env)
agent.train()
```

## Random Actions

```
total_rewards_random = 0
for _ in range(10000):
    env.reset()
    while not env.terminated:
        action = random.choice(env.actions)
        _, reward, _ = env.step(action)
        total_rewards_random += reward

print("Average total reward with random actions: ", total_rewards_random / 10000)
```

➡ Average total reward with random actions: 2.084

## With always spin slowly

```
# Evaluate the agent with always 'spin_slowly'
total_rewards_slowly = 0
for _ in range(10000):
    env.reset()
    while not env.terminated:
        action = 'spin_slowly'
        _, reward, _ = env.step(action)
        total_rewards_slowly += reward

print("Average total reward with 'spin_slowly': ", total_rewards_slowly / 10000)
```

➡ Average total reward with 'spin\_slowly': 3.0948

## With always spin rapidly

```
# Evaluate the agent with always 'spin_rapidly'
total_rewards_rapidly = 0
for _ in range(10000):
    env.reset()
```

```
while not env.terminated:
    action = 'spin_rapidly'
    _, reward, _ = env.step(action)
    total_rewards_rapidly += reward

print("Average total reward with 'spin_rapidly': ", total_rewards_rapidly / 10000)
```

➞ Average total reward with 'spin\_rapidly': 1.2872

```
total_rewards_random = 0
num_episodes = 10000
for _ in range(num_episodes):
    env.reset()
    while not env.terminated:
        action = random.choice(env.actions)
        next_state, reward, _ = env.step(action)
        print(f"Action: {action}, Reward: {reward}, Next State: {next_state}")
        total_rewards_random += reward

print("Average total reward with random actions: ", total_rewards_random / num_episodes)
```

➞ **Streaming output truncated to the last 5000 lines.**

Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 4, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 0, Next State: low  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 0, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 4, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 4, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 4, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 6, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 0, Next State: low  
Action: spin\_rapidly, Reward: 0, Next State: medium

Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 0, Next State: low  
Action: spin\_slowly, Reward: 0, Next State: low  
Action: spin\_slowly, Reward: 0, Next State: low  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 2, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 6, Next State: top  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: low  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: -2, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 6, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 8, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: low  
Action: spin\_rapidly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 2, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_slowly, Reward: 0, Next State: high  
Action: spin\_rapidly, Reward: 6, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: medium  
Action: spin\_rapidly, Reward: 0, Next State: high  
Action: spin\_slowly, Reward: 6, Next State: top  
Action: spin\_slowly, Reward: 0, Next State: medium

```
env = RoverEnv()  
agent = Agent(env)
```